

Small-data in the Big-data Era

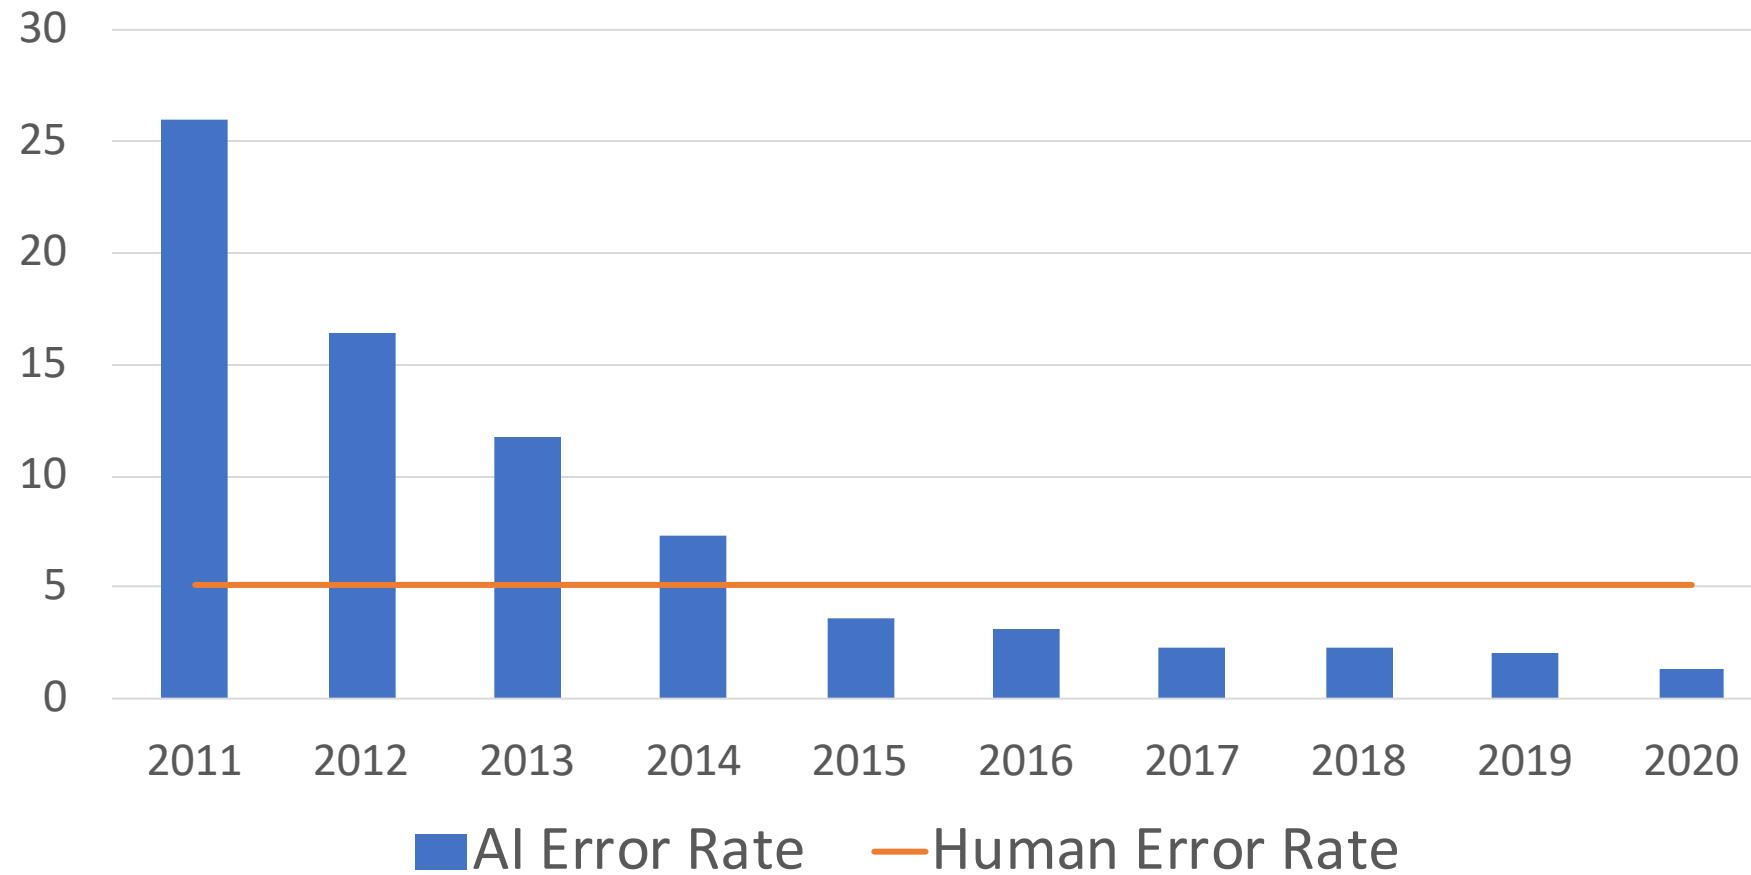
Eli Schwartz

Tel Aviv University; IBM Research AI

AI for Object Recognition



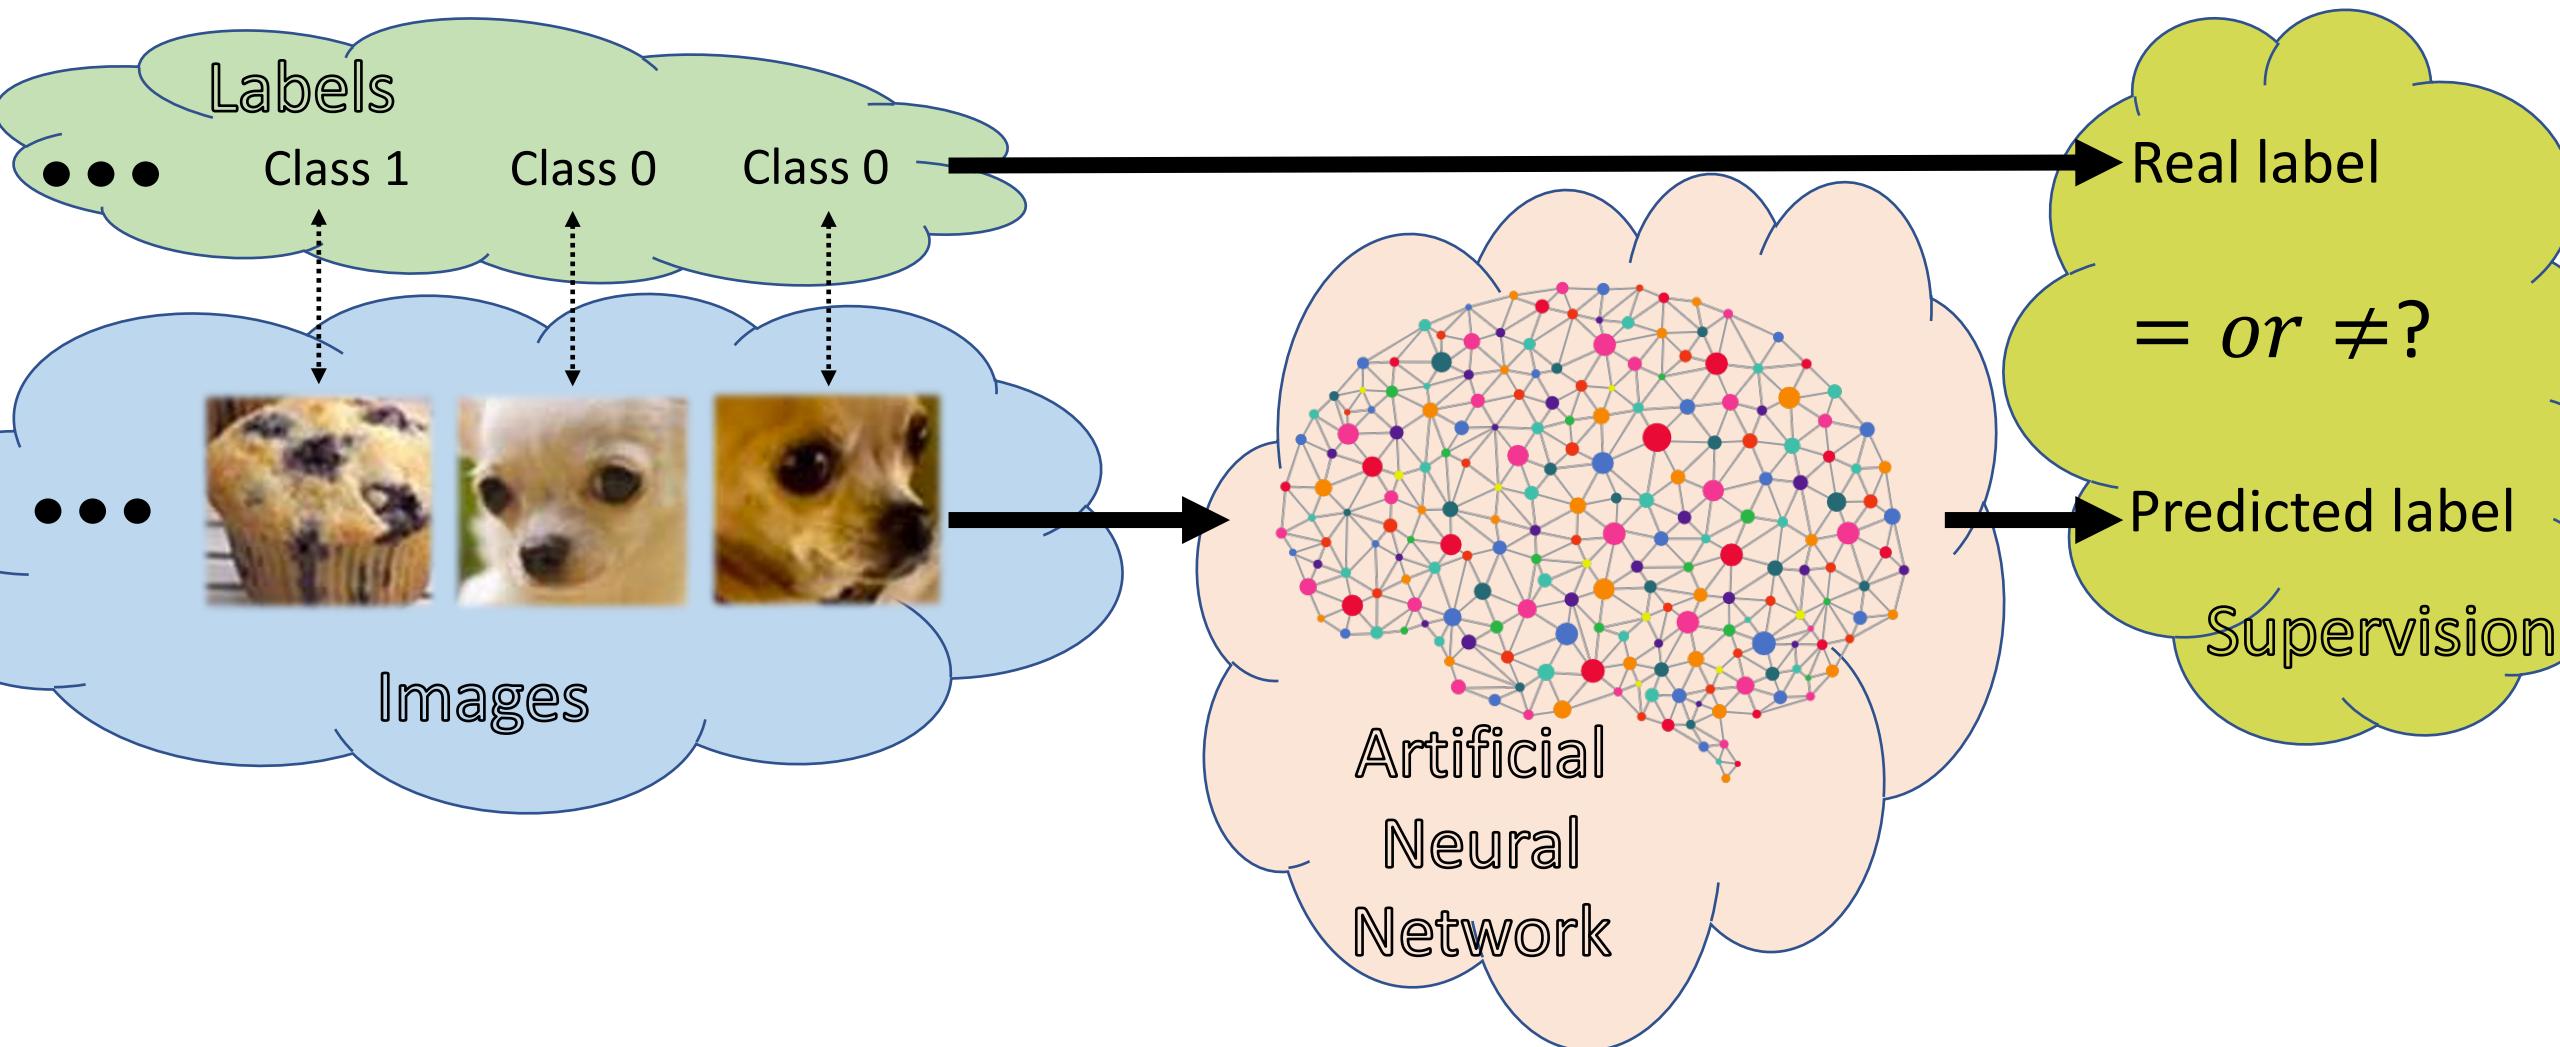
Object Recognition Performance



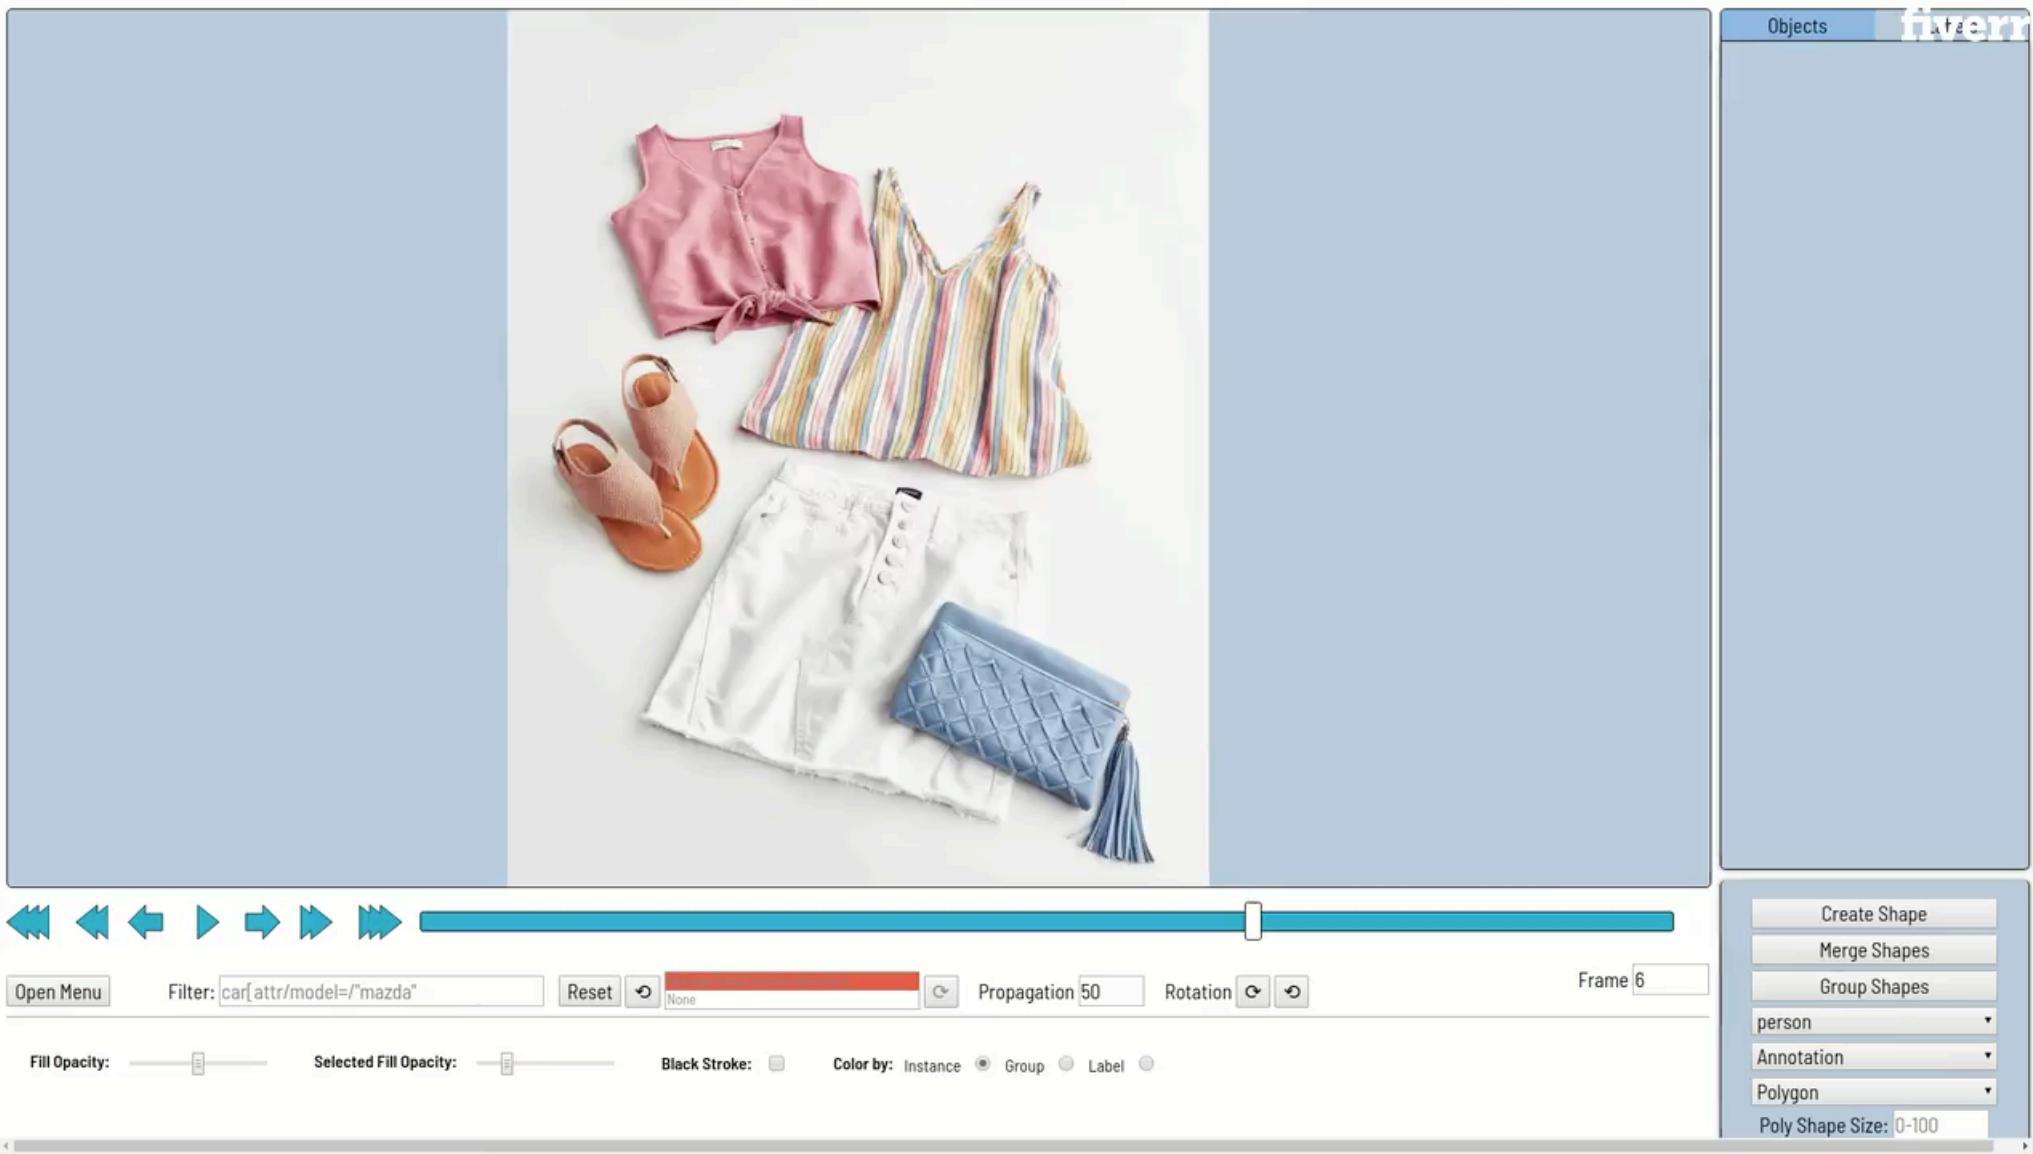
Big-Data



Supervised Learning 101



Labeling is Hard



Labeling is Hard and costly



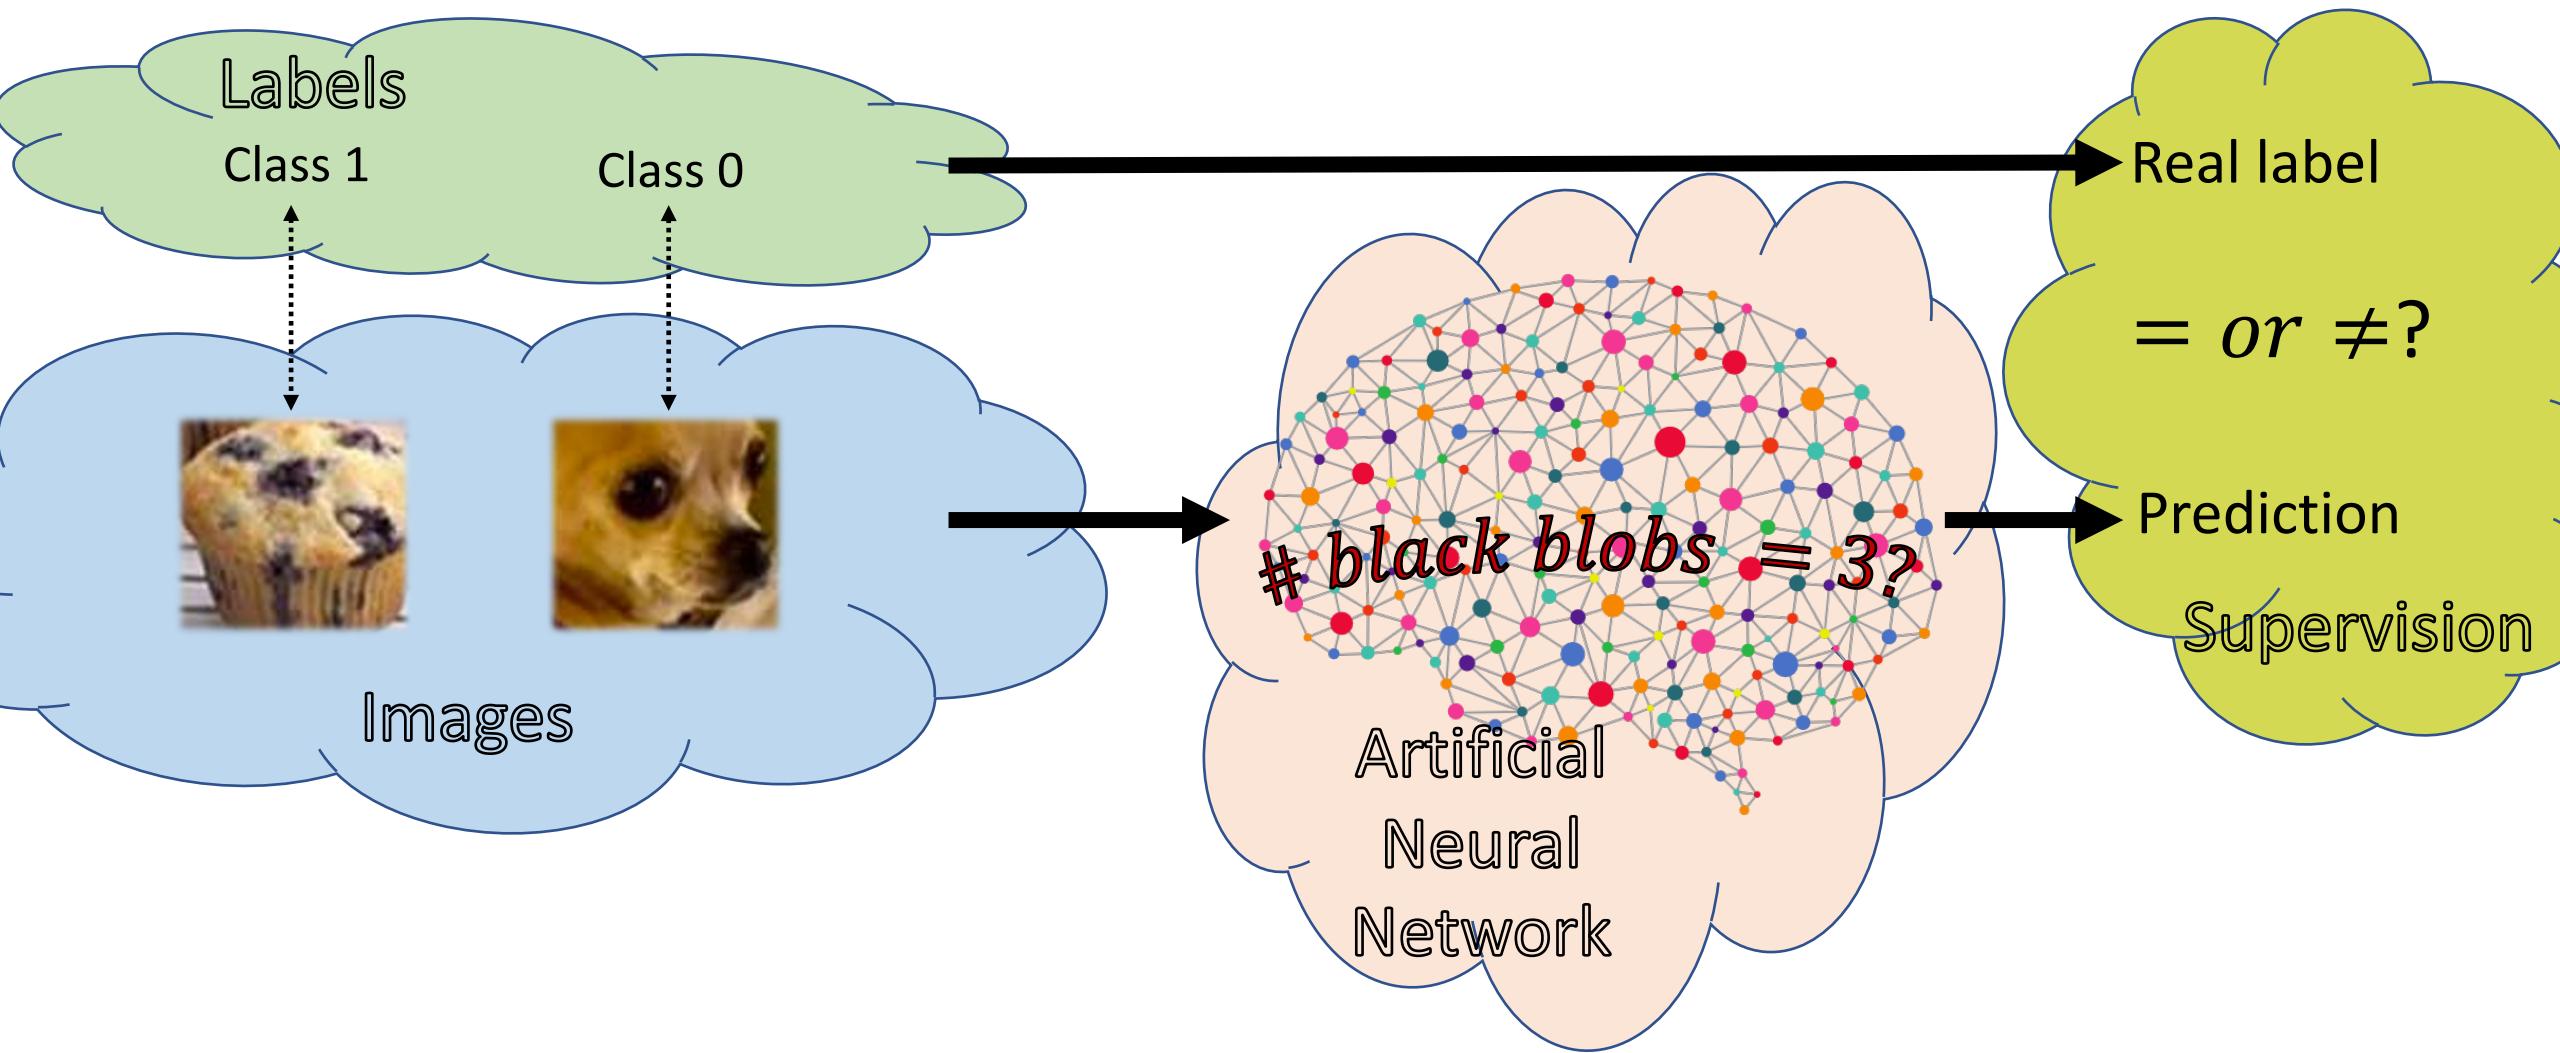
Humans need much less supervision



Humans need much less supervision

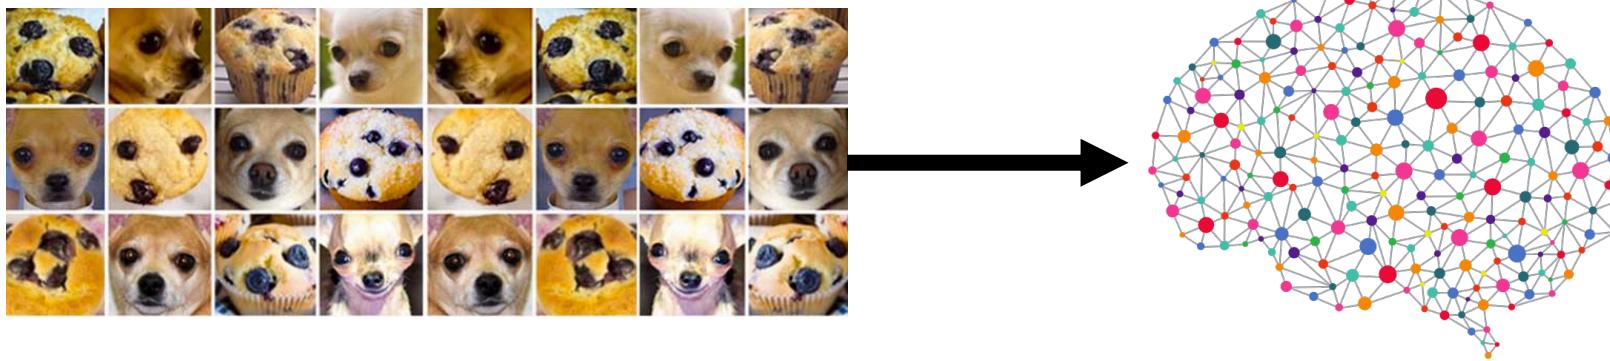


Standard training breaks for small-data



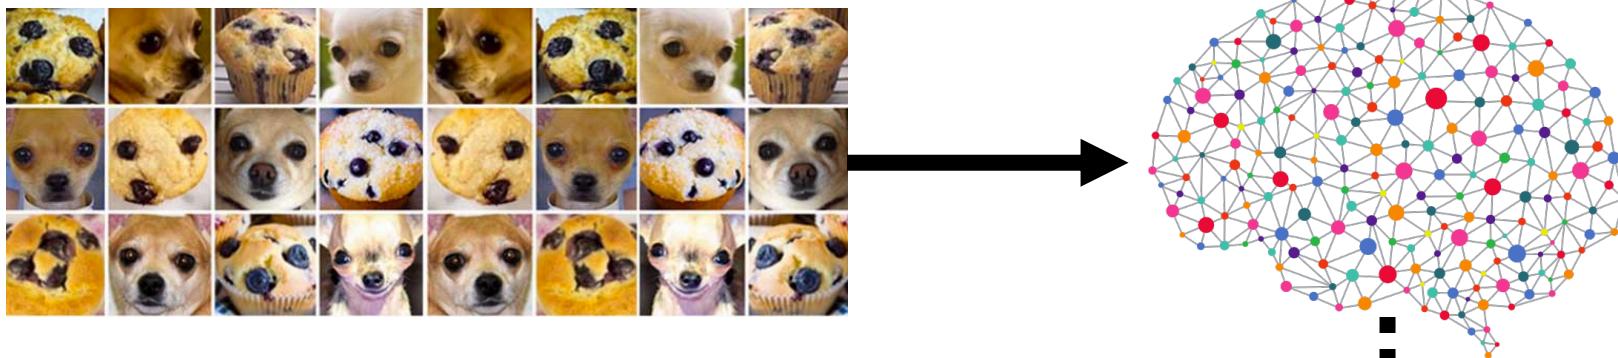
One-Shot Learning

Phase 1 – Training:

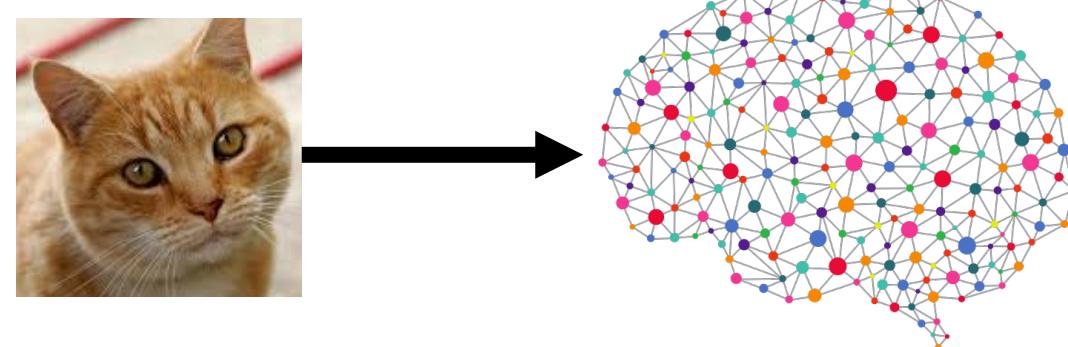


One-Shot Learning

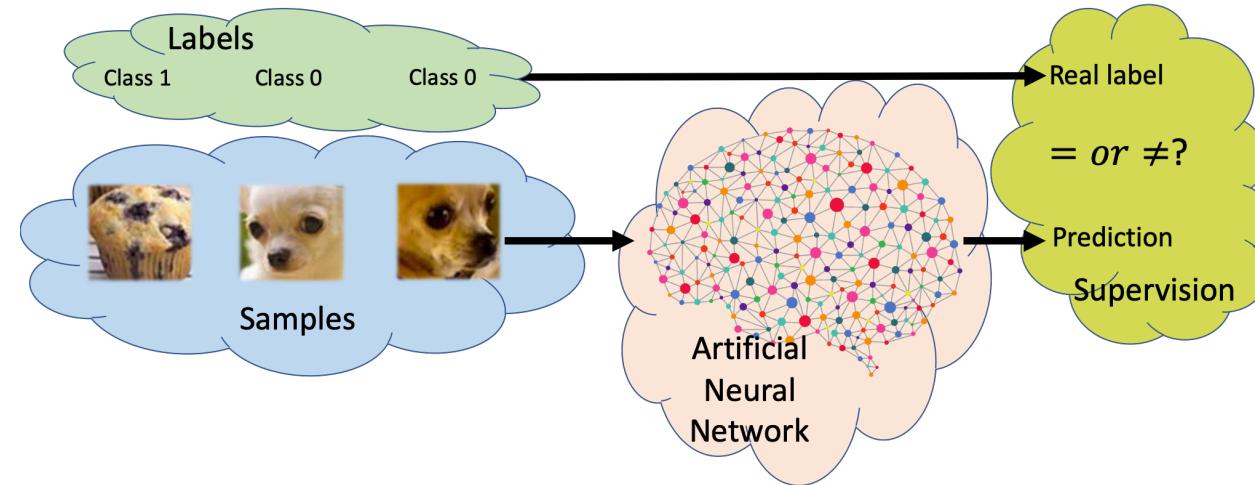
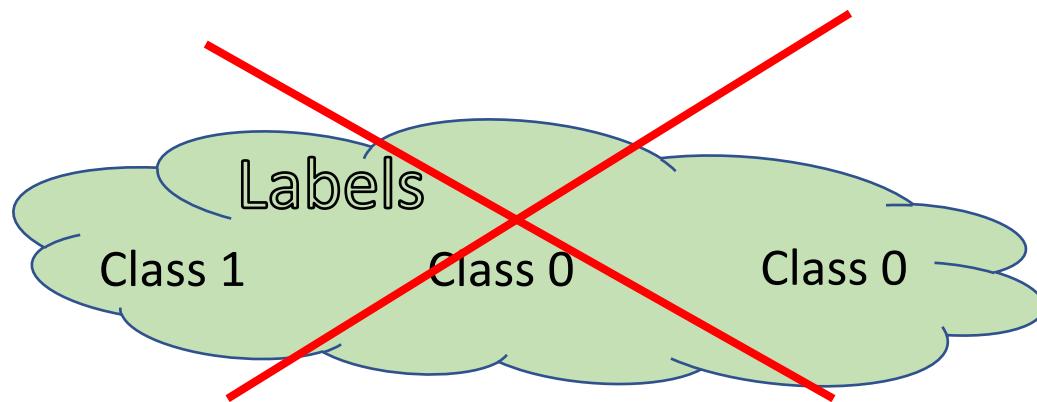
Phase 1 – Training:



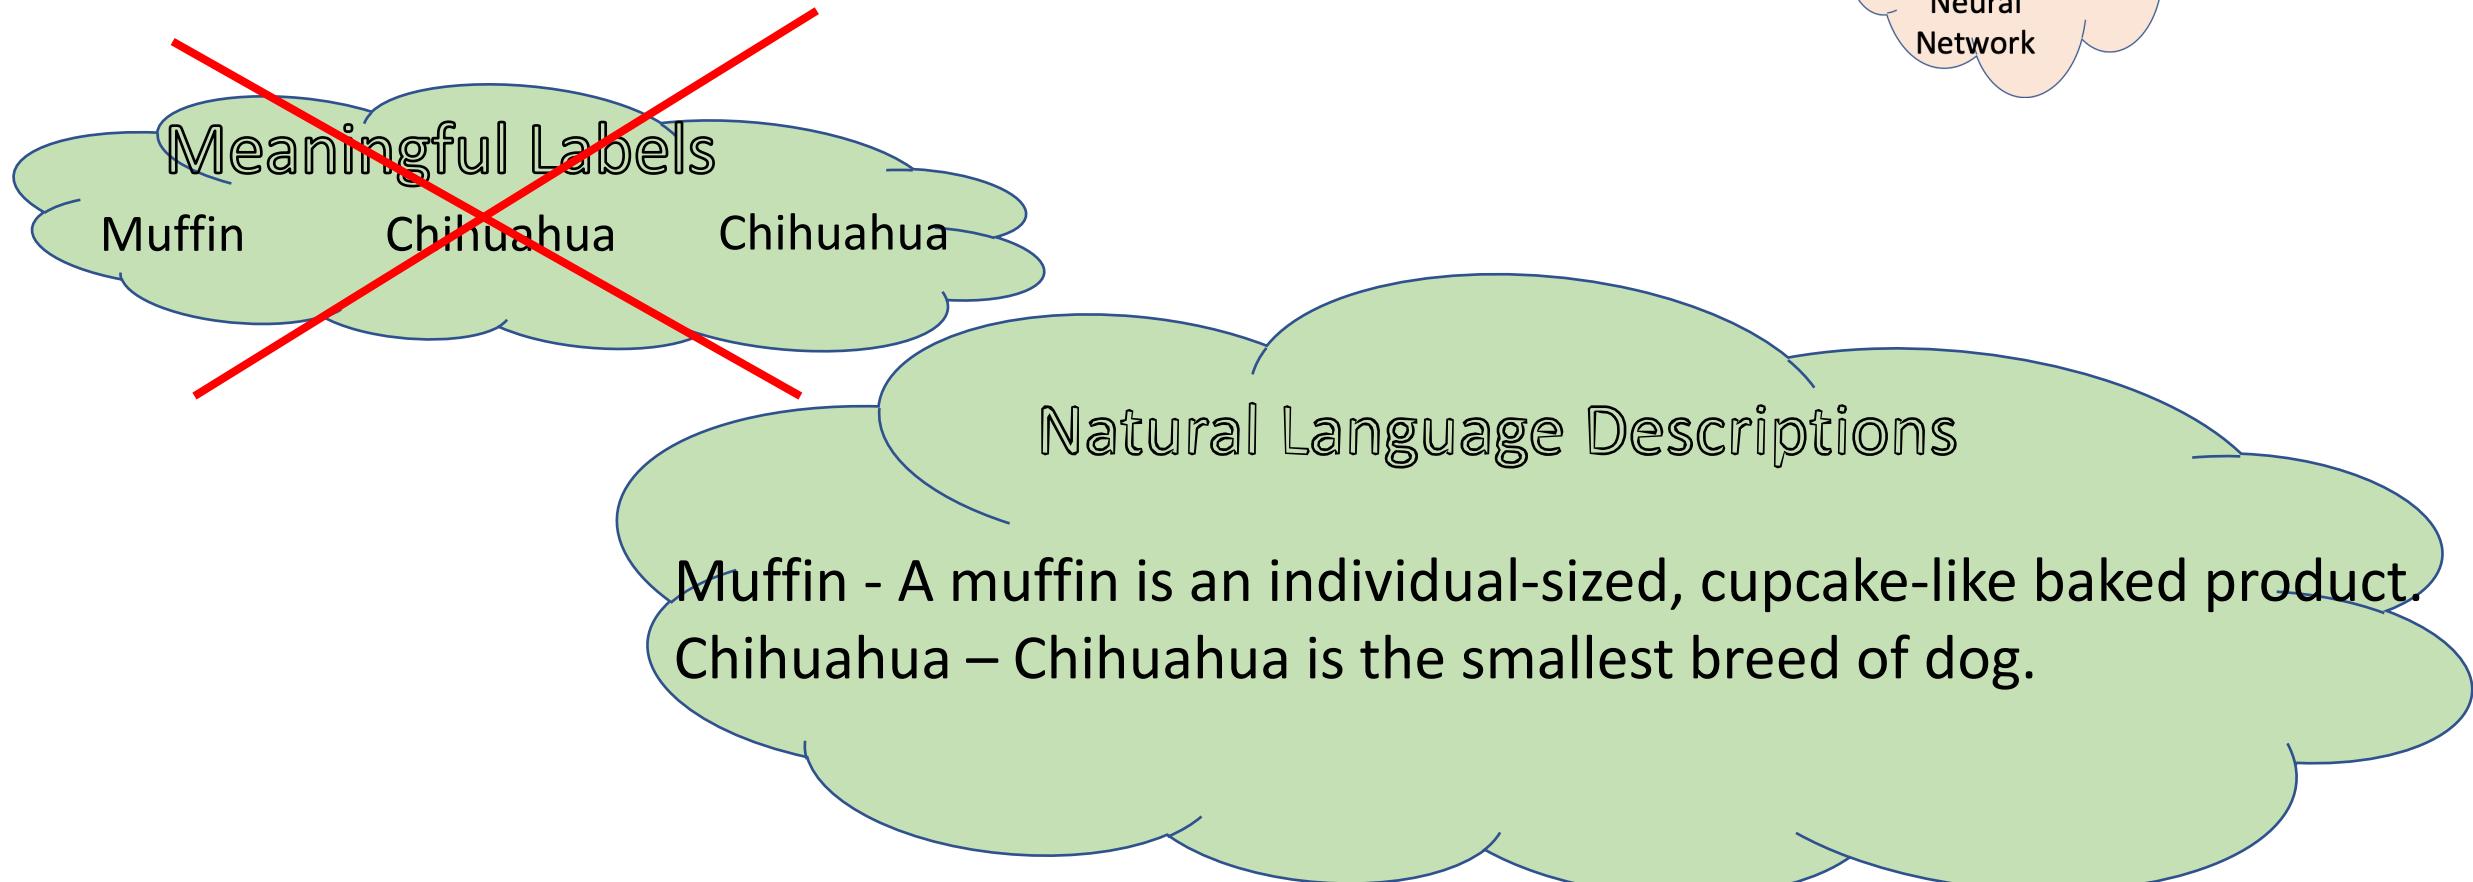
Phase 2 – Adapting to novel classes:



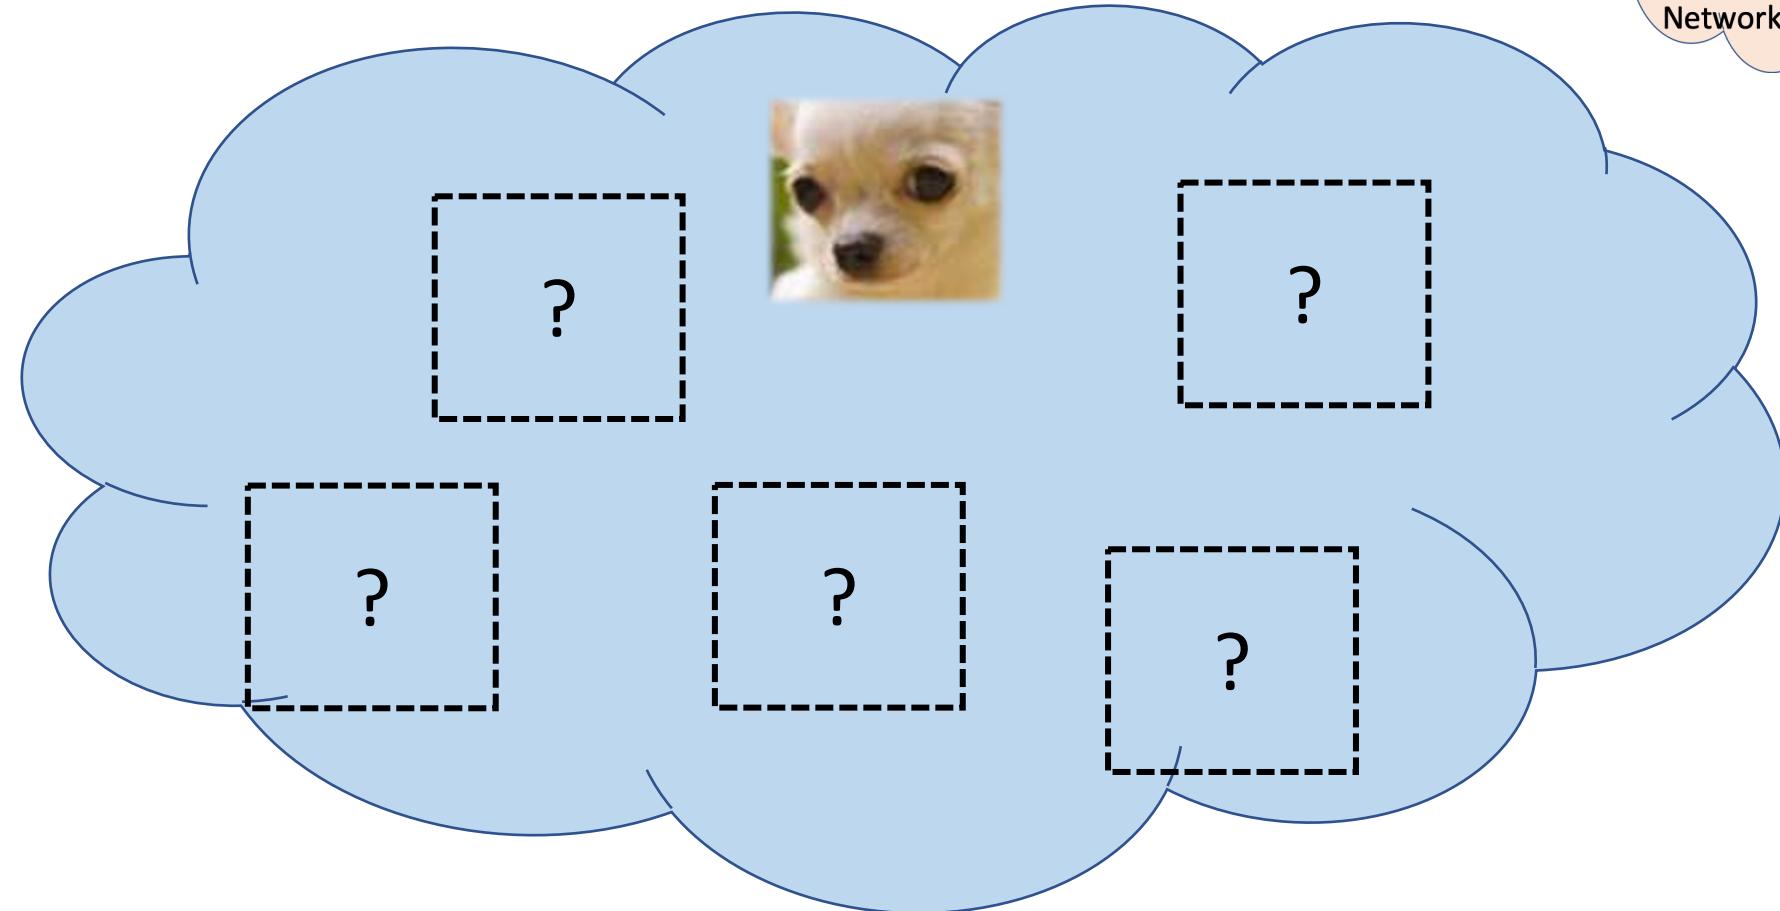
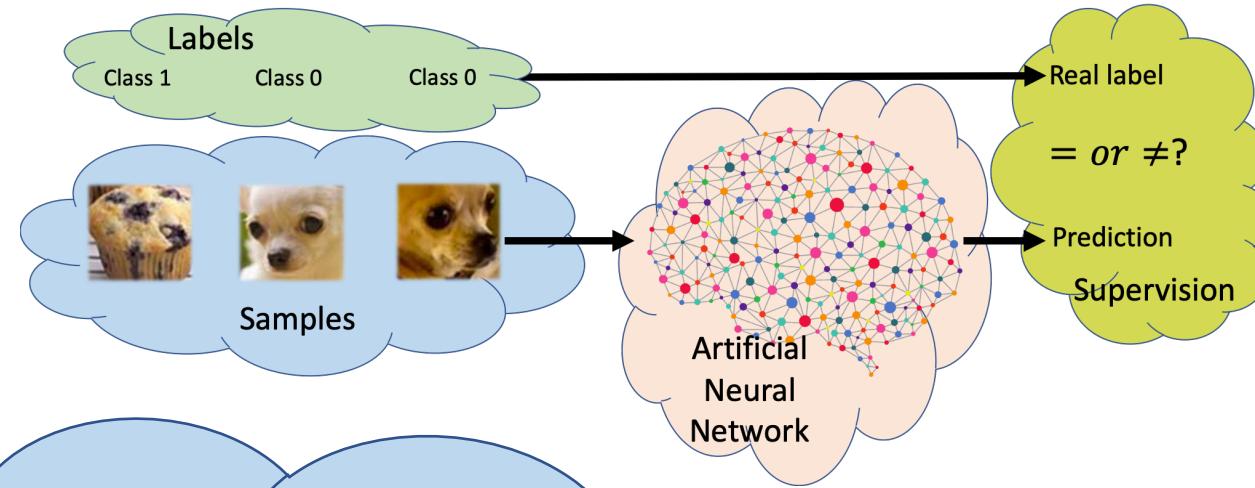
Improving the labels



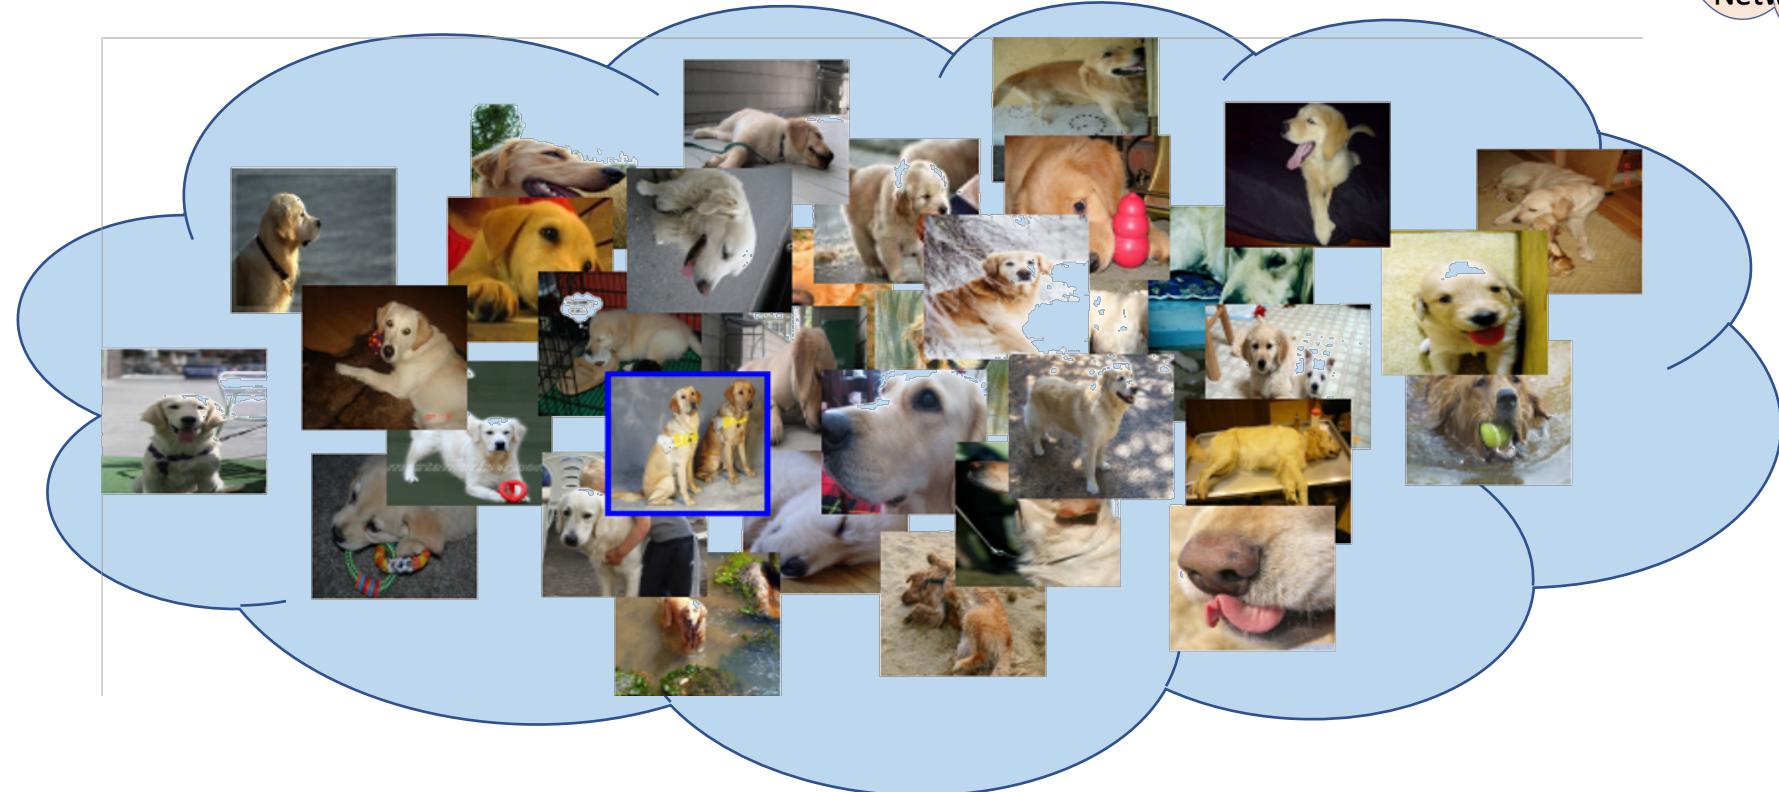
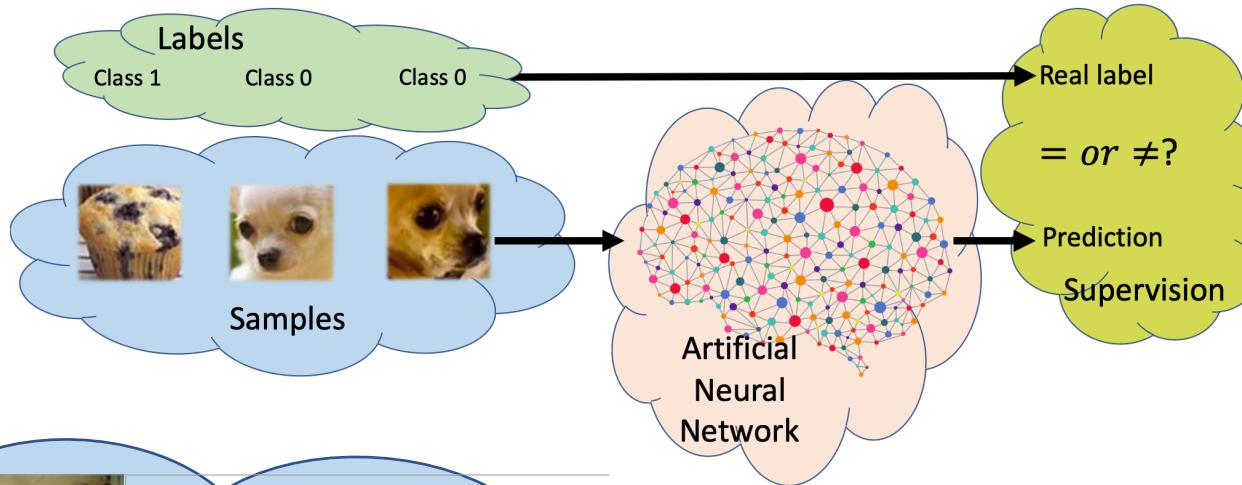
Improving the labels



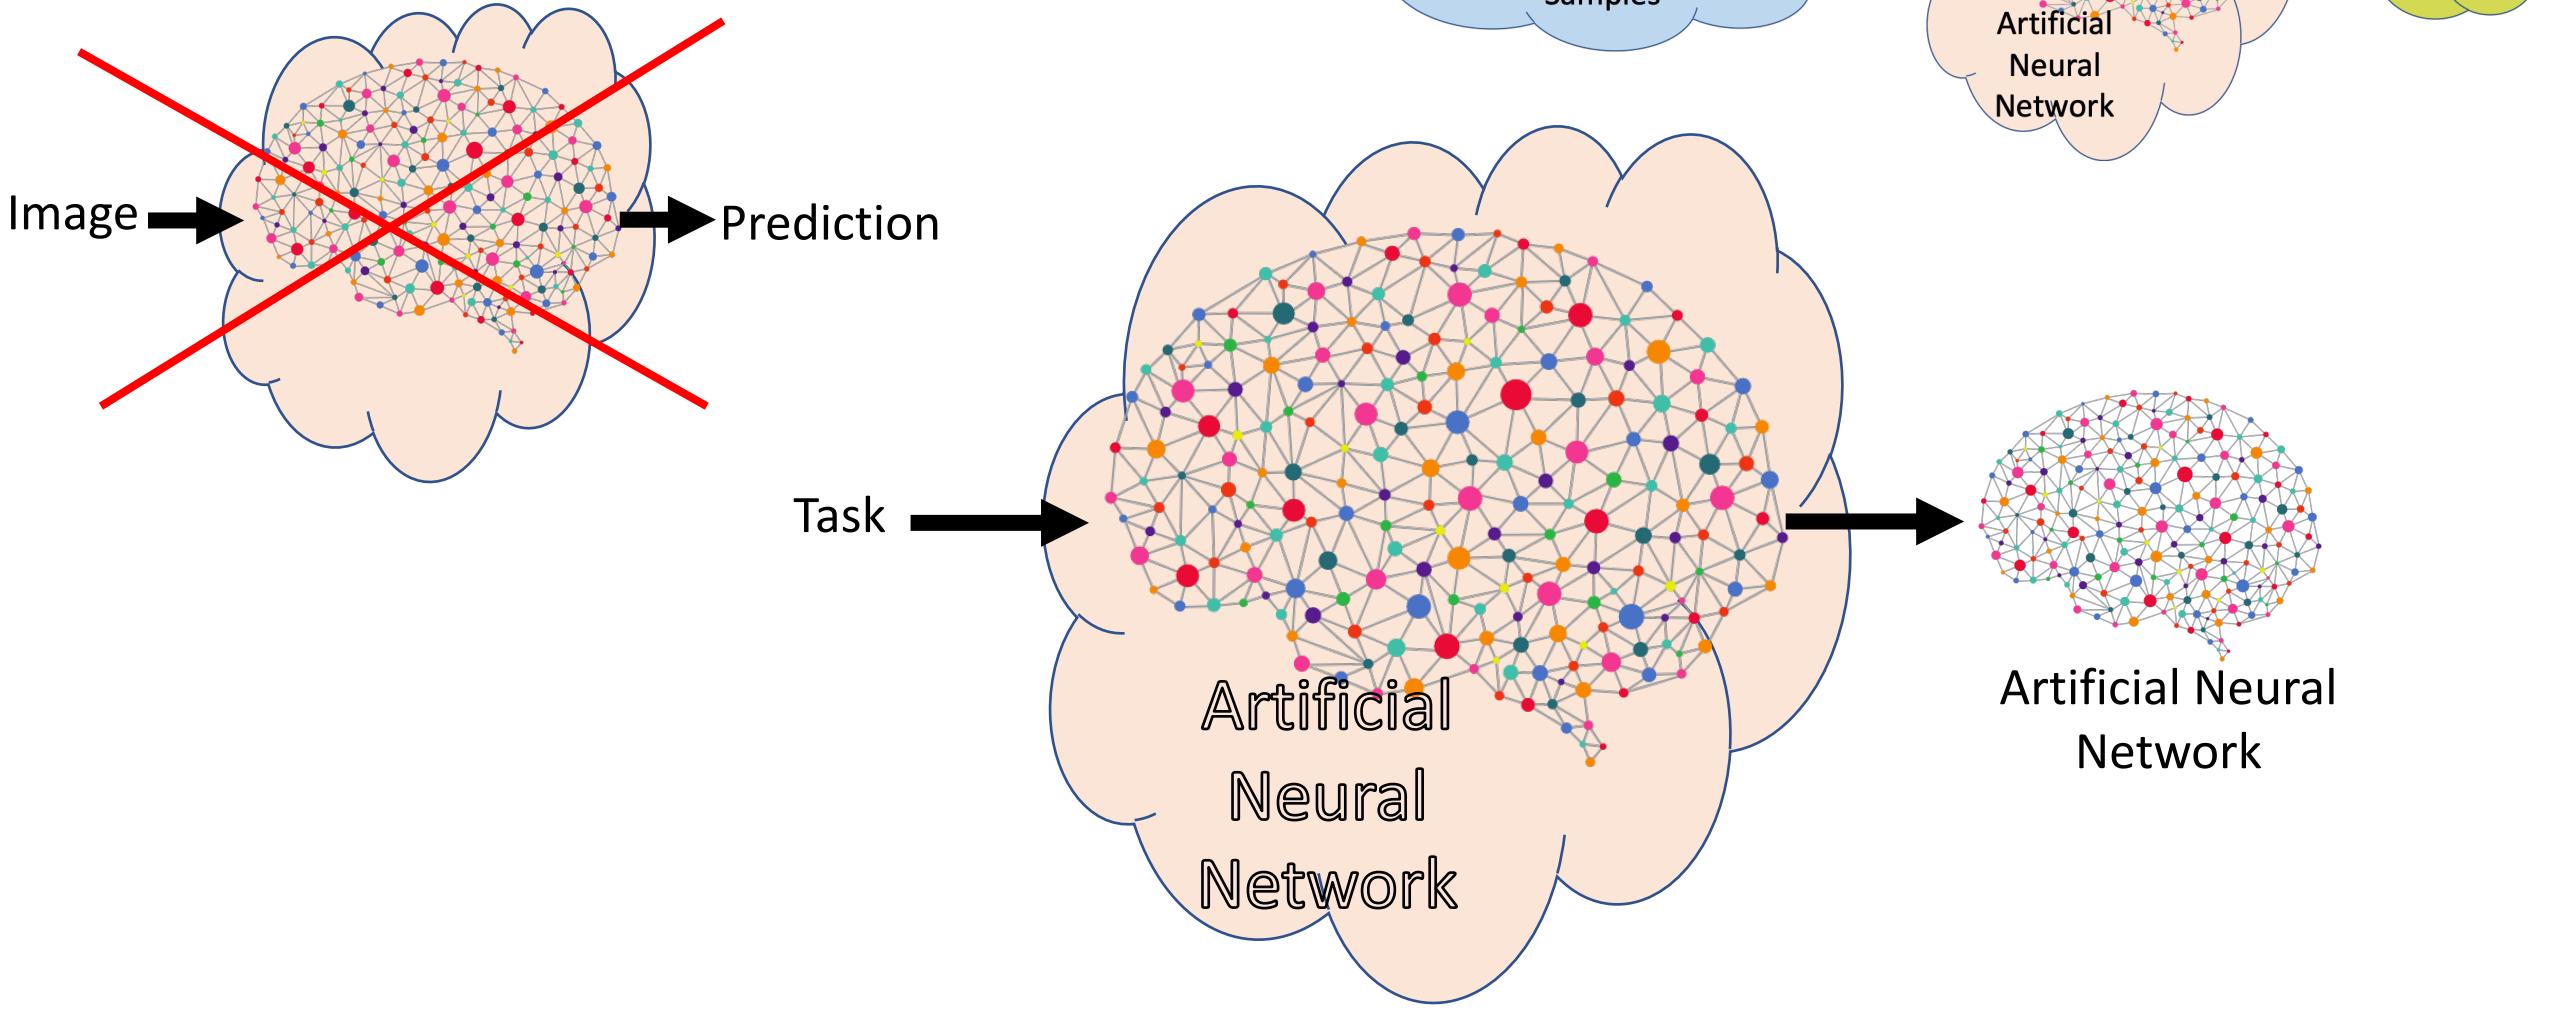
Hallucinating More Images



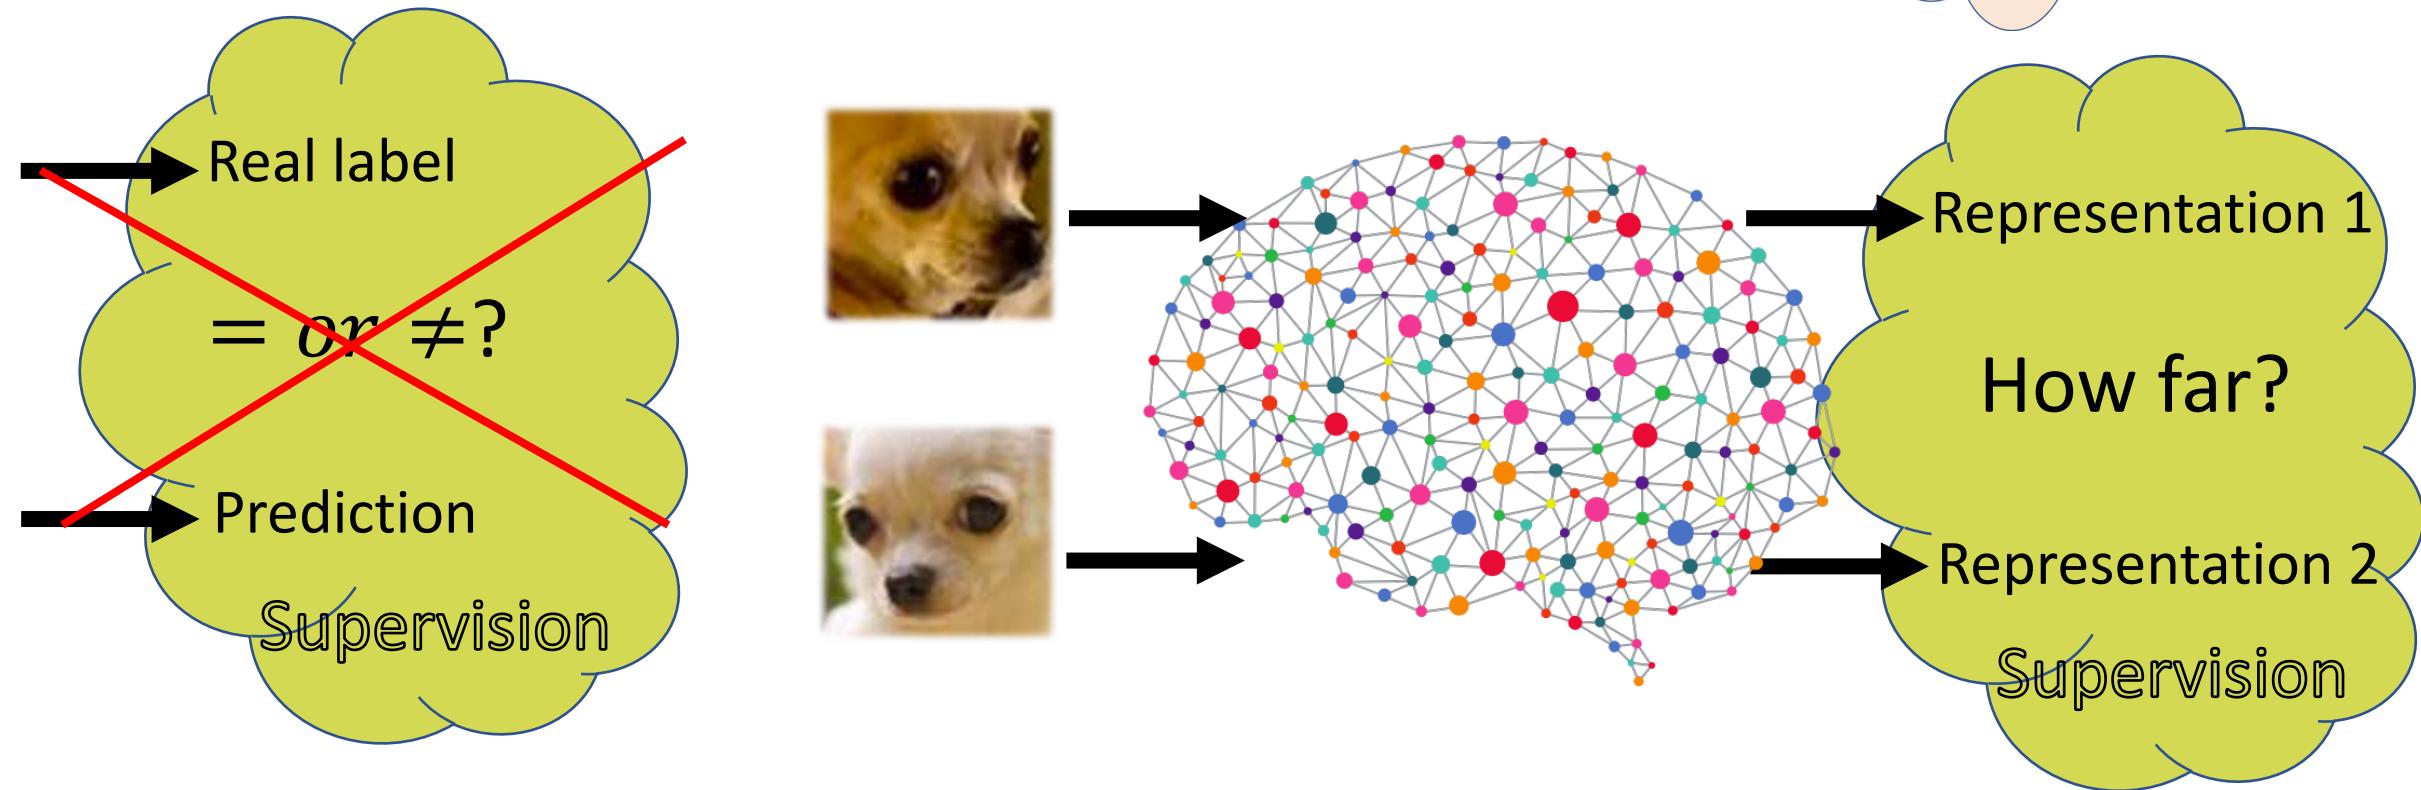
Hallucinating More Images



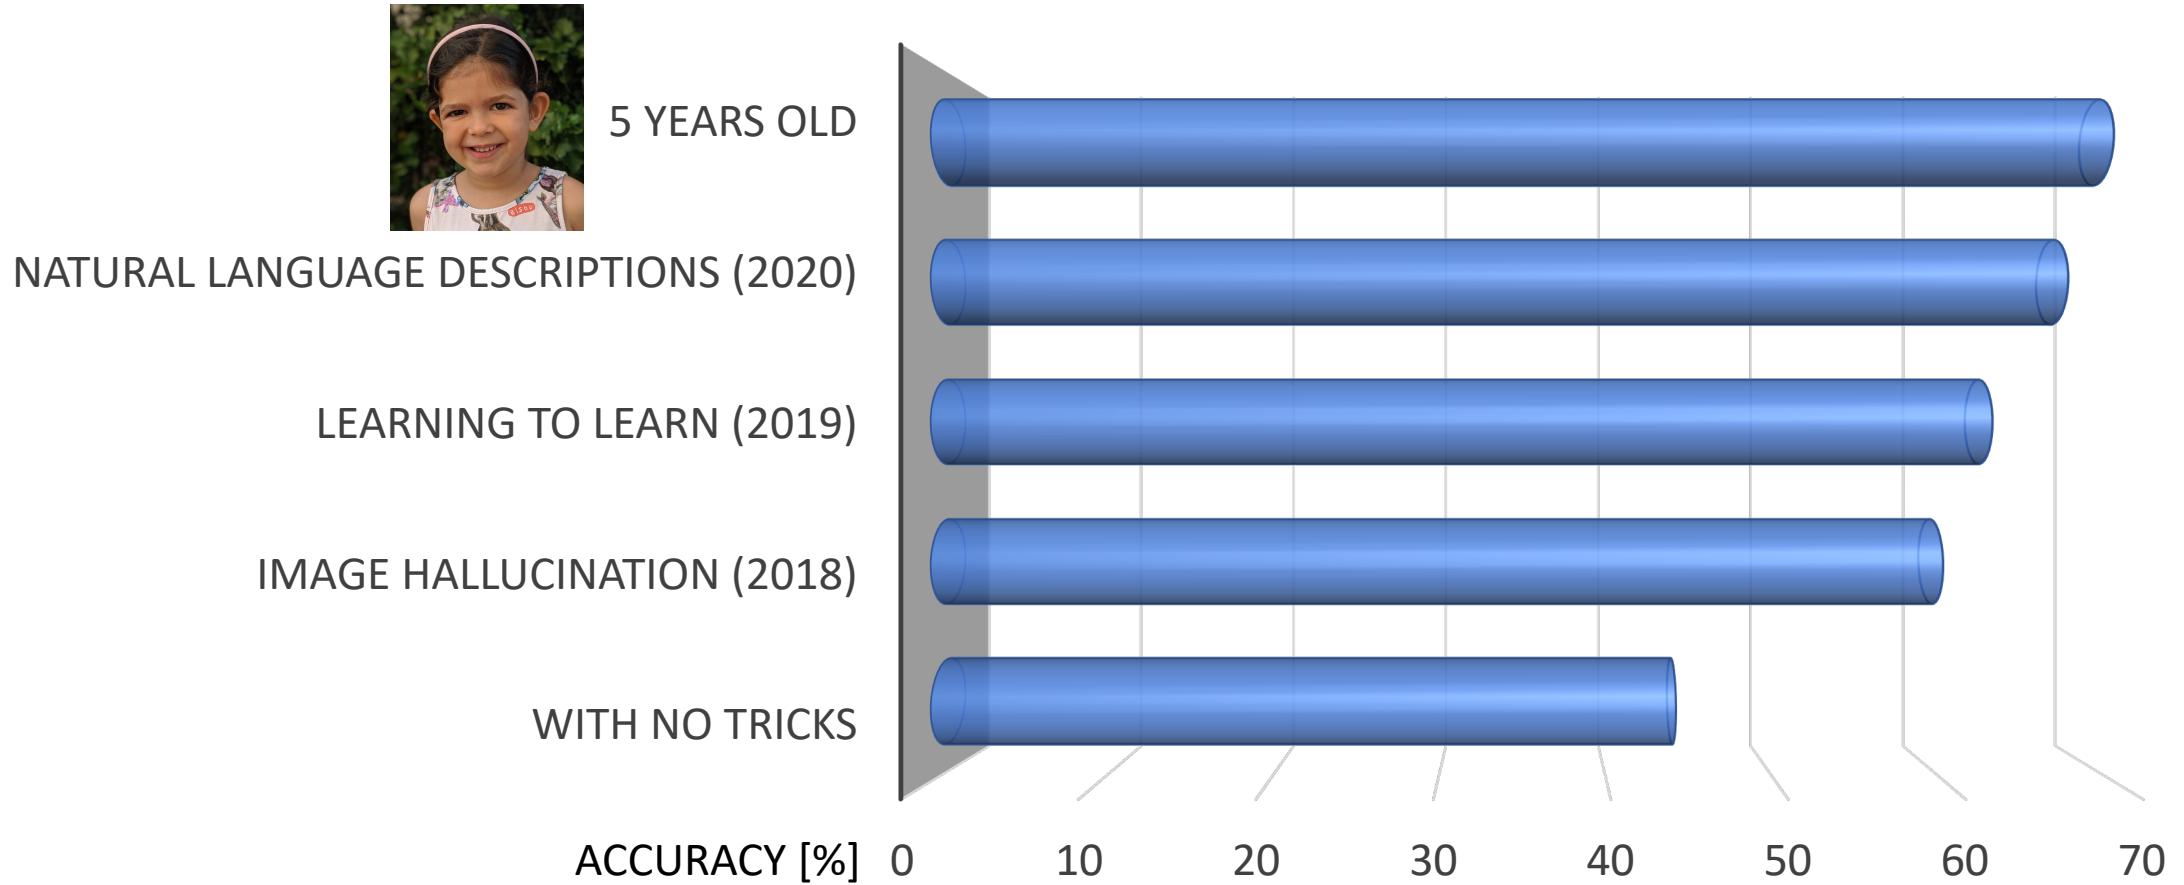
“Meta-Learning” or “Learning to Learn”

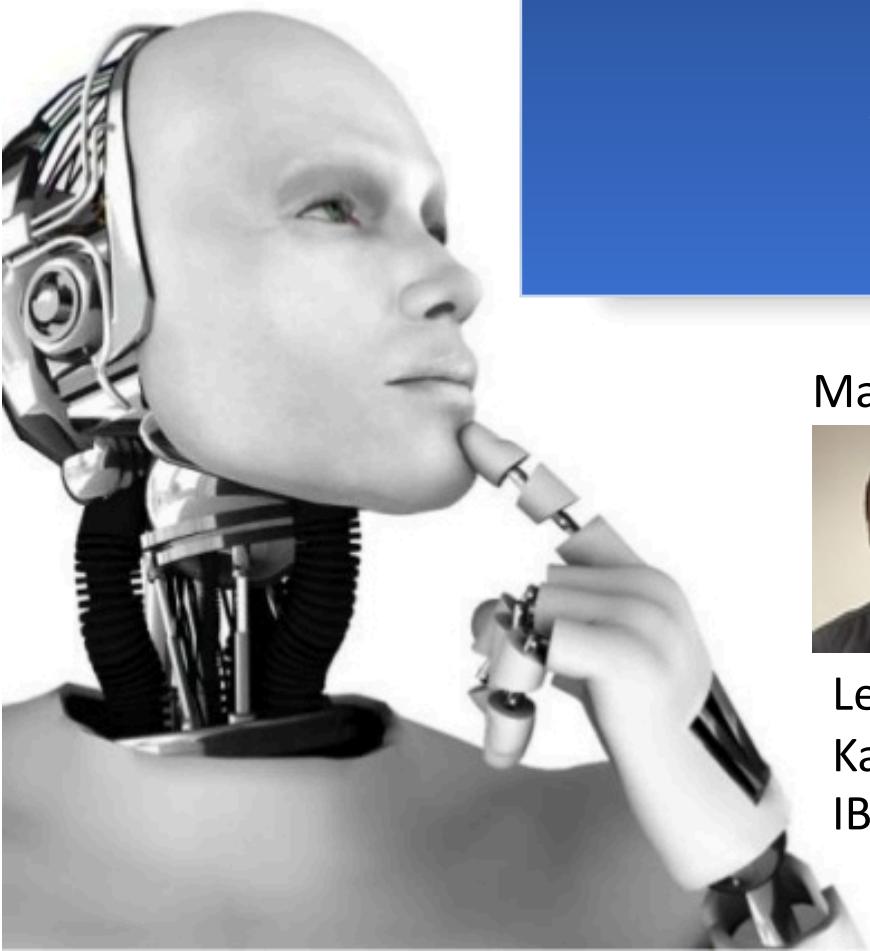


Asking the Right Question



Learning from a single example performance





Questions?

Thank you!

Main collaborators:



Leonid
Karlinsky,
IBM



Joseph
Shtok,
IBM



Sivan
Harary,
IBM



Sivan
Doveh,
IBM



Rogerio
Ferris,
IBM



Alex
Bronstein,
Technion



Raja
Giryes,
TAU